


Introduction to Operating System

Alok Jhaldiyal
Assistant Professor
SoCS, UPES
Dehradun

General Information

- Textbook:
 - Operating System Concepts, 8th or 9th Ed, by Silberschatz, Galvin, and Gagne
- Reference Books:
 - Operating Systems: Principles and Practice by Anderson and Dahlin
 - Modern Operating Systems by Andrew Tanenbaum
- Programming assignments will be covered in associated Lab.

Our Exposure to various OS's



Computer System

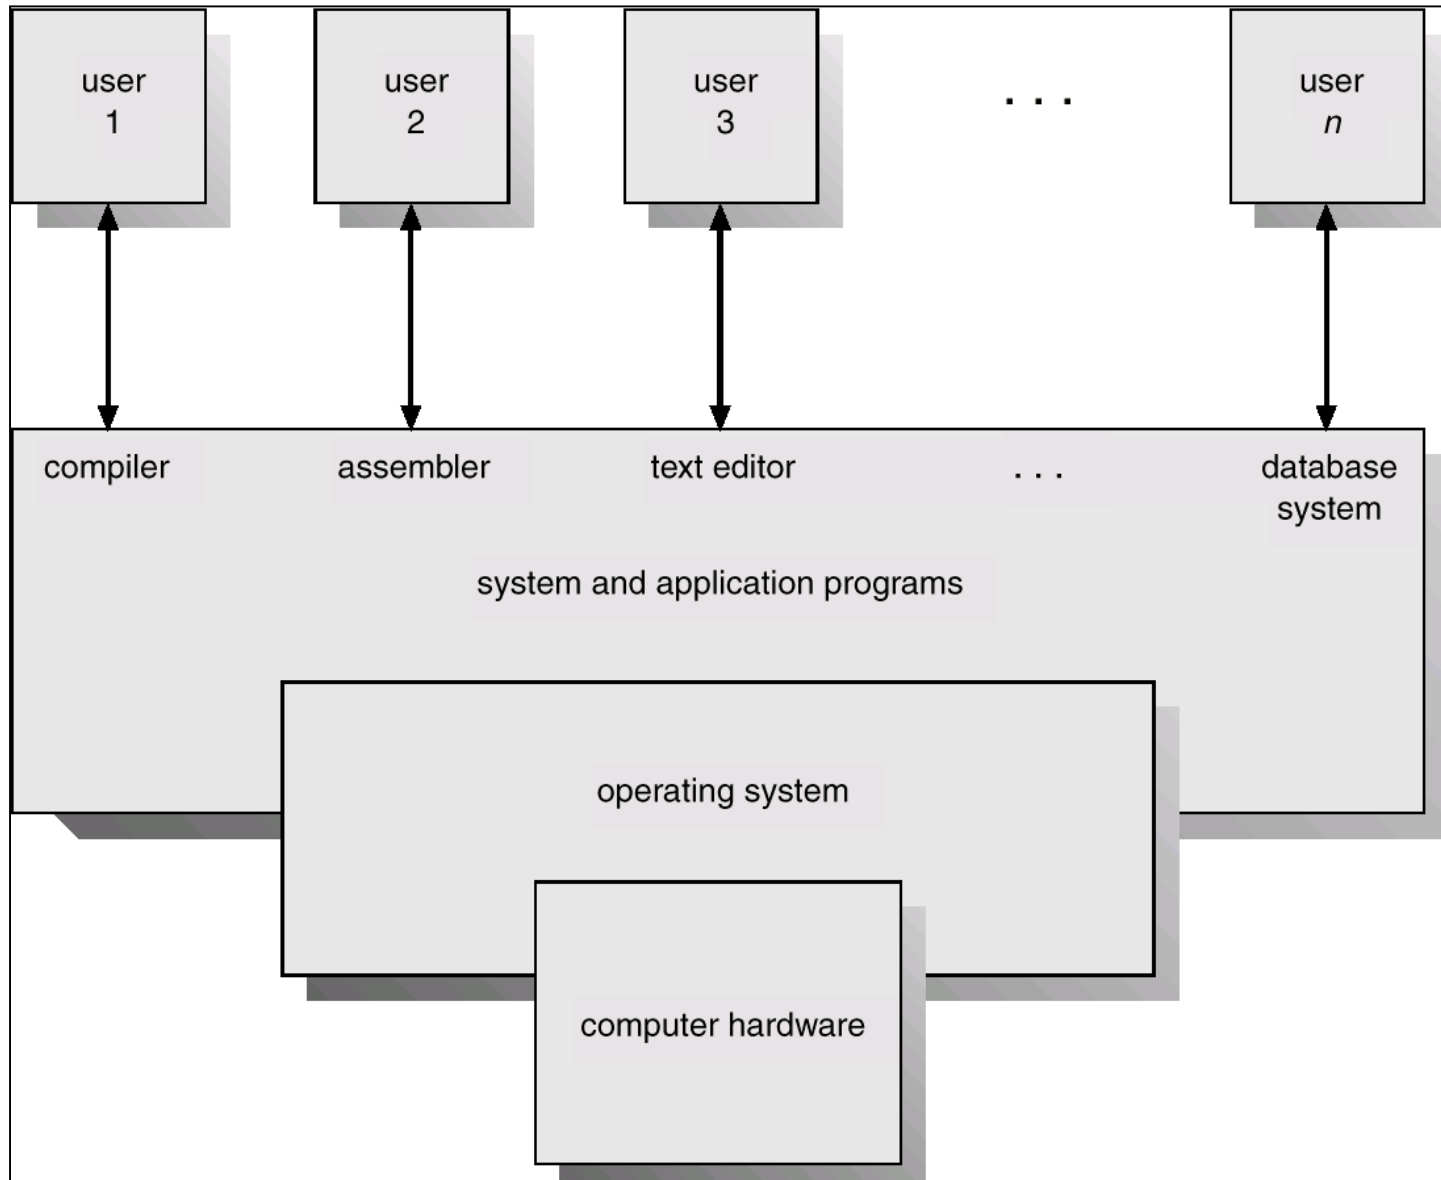
- computer system can be divided roughly into four components:
 - **Hardware**
 - The **central processing unit (CPU)**, the **memory**, and the **input/output (I/O) devices**
 - Provides the basic computing resources for the system
 - **Operating System**
 - Controls the hardware and coordinates its use among the various application programs for the various users
 - Operating System provides the means for proper use of these resources.
 - **Application Programs**
 - Define the ways in which these resources are used to solve users' computing problems
 - Application Programs: Word Processors, spreadsheets, compilers, browsers.....
 - **Users**

What is operating system?

- User-centric definition
 - A program that acts as an intermediary between a user of a computer and the computer hardware
 - Defines an interface for the user to use services provided by the system
 - Provides a “view” of the system to the user Converts what the hardware gives to what the user wants
 - The view can hide many details of the hardware that the user does not need to know
 - Can even give a very different view of the operating environment to the user than what is actually there

- System-centric definition
 - Efficiently manages and allocates resources to users
 - Controls the execution of user programs and operations of I/O devices
 - Provides isolation/protection between different user programs

Abstract View of System Components



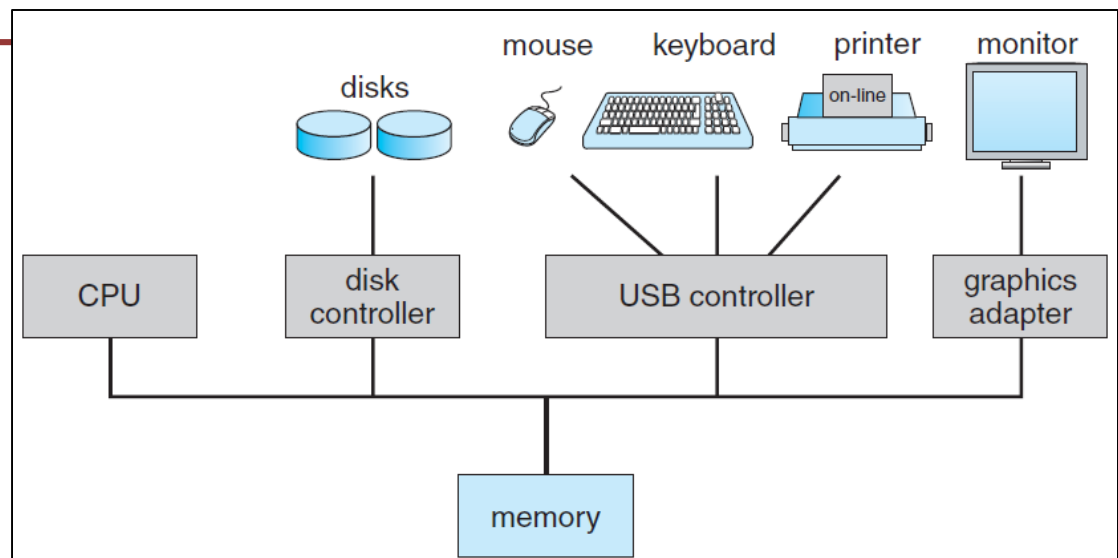
Types of System

- Batch Systems
 - Multiple jobs, but only one job in memory at one time and executed (till completion) before the next one starts
- Multiprogrammed Batch Systems
 - Multiple jobs in memory, CPU is multiplexed between them
 - CPU-bound vs I/O bound jobs
- Time-sharing Systems
 - Multiple jobs in memory and on disk, CPU is multiplexed among jobs in memory, jobs swapped between disk and memory
 - Allows interaction with users

- Personal Computers
 - Dedicated to a single user at one time
- Multiprocessing Systems
 - More than one CPU in a single machine to allocate jobs to
 - Symmetric Multiprocessing, NUMA machines ...
 - Multicore
- Other Parallel Systems, Distributed Systems, Clusters...
 - Different types of systems with multiple CPUs/Machines
- Real Time Systems
 - Systems to run jobs with time guarantees
- Many other types
 - Embedded systems, mobiles/smartphones,

Computer System Architecture

- Computer system consists of one or more CPUs and a number of device controllers, which execute in parallel.
- Each device controller is in charge of a specific type of device.
- For orderly access to the shared memory, a memory controller synchronizes access to the memory



- Bootstrap program
 - Once computer is powered on, to start running a initial program: Bootstrap Program
 - Initializes all aspects of the system, from CPU registers to device controllers to memory contents.
 - Stored within the computer hardware in read-only memory (**ROM**)
 - Bootstrap program locates the operating-system kernel and loads it into memory
 - A loaded and executing kernel provides services of the system to the user.
- Some services are provided outside of the kernel, by system programs that are loaded into memory at boot time to become **system processes**

Bootstrap Steps

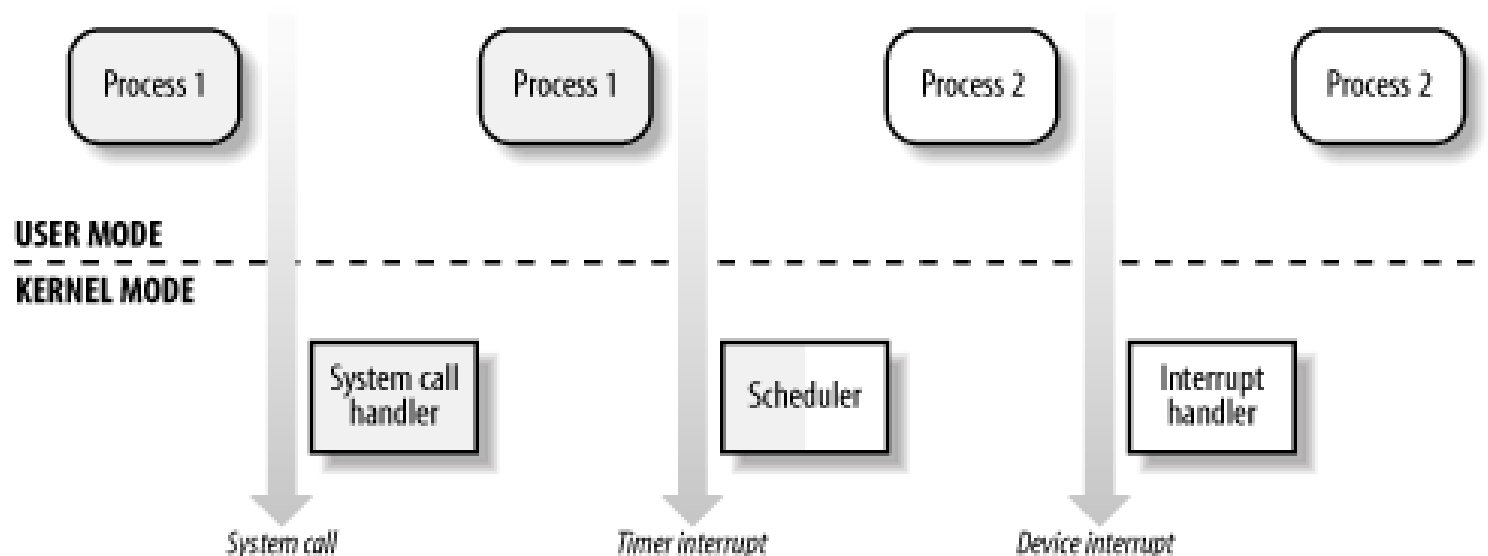
- Loading the code and initializing the kernel
- Detecting the Devices and configuring them
- Creating spontaneous system processes
- Operator intervention (manual boot only)
- Execution of system startup scripts
- Multiuser operation

Kernel

- Core component of operating system
- Unix kernels provide an execution environment in which applications may run
- Interfaces between the three major computer hardware components
 - Application/user interface
 - The CPU
 - Memory
 - I/O devices.
- Kernel also sets up memory address space for applications, loads files with application code into memory

Transitions Between User and Kernel Mode

- Process 1 in User Mode issues a system call, after which the process switches to Kernel Mode and the system call is serviced.
- Process 1 then resumes execution in User Mode until a timer interrupt occurs and the scheduler is activated in Kernel Mode
- A process switch takes place and Process 2 starts its execution in User Mode until a hardware device raises an interrupt.
- As a consequence of the interrupt, Process 2 switches to Kernel Mode and services the interrupt.



Homework

Case Study 1: Study and document the Unix bootstrap process.

Submission Date: 18/08/2023

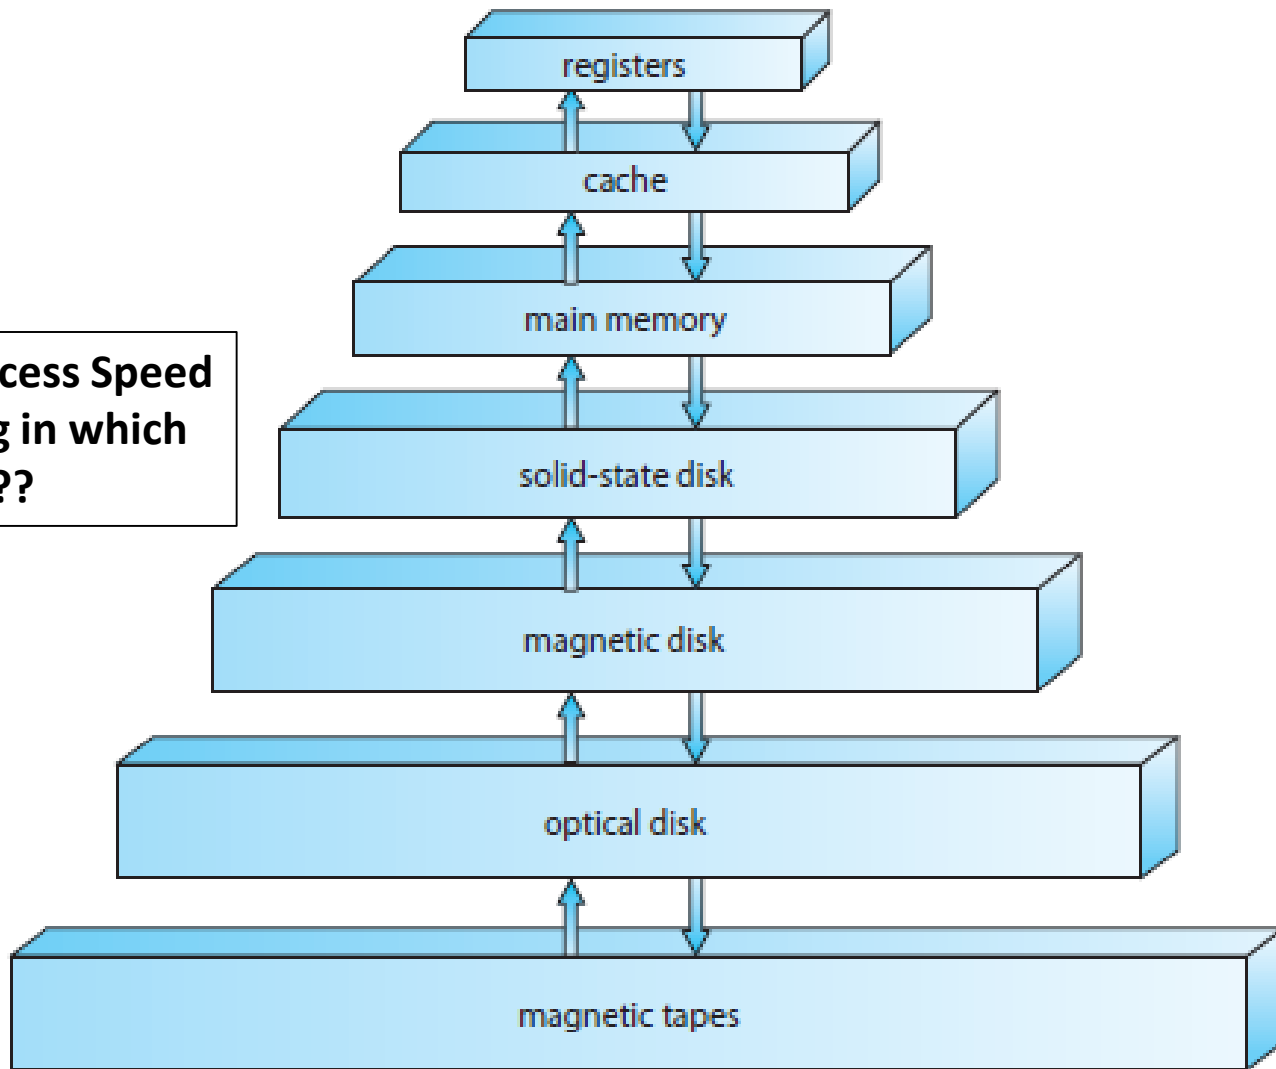
- Interrupts

- Occurrence of an event is usually signalled by an **interrupt** from either the hardware or the software.
- When the CPU is interrupted, it stops what it is doing and immediately transfers execution to a fixed location.
- Fixed location usually contains the starting address where the service routine for the interrupt is located.
- On completion of service routine the CPU resumes the interrupted computation.

Storage Structure

- CPU can load instructions only from memory, so any programs to run must be stored there.
- Mostly computer access there programs from Random Access Memory (RAM)
- Static programs such as Bootstrap are stored in ROM.
- All forms of memory provide an array of bytes, each byte having its own address.
- Sequence of load or store instructions maintains the execution.
- Load instruction moves a byte or word from main memory to an internal register within the CPU.
- Store instruction moves the content of a register to main memory

**Memory Access Speed
is increasing in which
direction????**



Storage-device hierarchy

I/O Structure


- A large portion of operating system code is dedicated to managing I/O.
- Device controller who is in charge of a specific type of device, maintains some local buffer storage and a set of special-purpose registers.
- To start an I/O operation, the device driver loads the appropriate registers within the device controller.
- Device driver then informs the operating system that the data is loaded and is ready.

Operating System Operations

- Process Management
 - A program in execution, as mentioned, is a process.
 - A word-processing program being run by an individual user on a PC is a process.
 - A process needs certain resources—CPU time, memory, files, and I/O devices.
 - Activities in connection with process management:
 - Scheduling processes and threads on the CPUs
 - Creating and deleting both user and system processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization

- **Memory Management**

- Main memory is central to the operation of a modern computer system
- Main memory is a repository of quickly accessible data shared by the CPU and I/O devices.
- For a program to be executed, it must be mapped to absolute addresses and loaded into memory.
- Program executes by accessing program instructions from memory by their absolute addresses.
- Program terminates and its memory space is declared available
- Activities in connection with memory management:
 - Keeping track of which parts of memory are currently being used and who is using them
 - Deciding which processes (or parts of processes) and data to move into and out of memory.
 - Allocating and deallocating memory space as needed.



- Absolute Address:
 - A fixed address in memory
 - Is not a relative address
 - Also called real addresses

- Storage Management
 - Operating system provides a uniform, logical view of information storage, a logical unit, the **file**
 - OS maps files onto physical media and accesses these files via the storage devices.
 - Activities in connection with storage management:
 - File System Management
 - Mass storage management
 - Caching
 - I/O Systems
 - File system management:
 - Creating and deleting files
 - Creating and deleting directories to organize files
 - Supporting primitives for manipulating files and directories
 - Mapping files onto secondary storage
 - Backing up files on stable (nonvolatile) storage media

- Storage Management
 - Mass-Storage Management
 - Free-space management
 - Storage allocation
 - Disk scheduling
 - Caching
 - I/O Systems
 - A memory-management component that includes buffering, caching, and spooling
 - A general device-driver interface
 - Drivers for specific hardware devices


Performance of Various Levels of Storage

- Movement between levels of storage hierarchy can be explicit or implicit

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape

Migration of Integer A from Disk to Register

- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide cache coherency in hardware such that all CPUs have the most recent value in their cache

- Protection and Security

- Protection, then, is any mechanism for controlling the access of processes or users to the resources defined by a computer system.
- Job of **security** is to defend a system from external and internal attacks which may include viruses and worms, denial-of service attacks, identity theft and theft of service.

Modes of Operating System

- Operating System is interface between applications and the underlying hardware.
- At the same time to maintain systems integrity, it needs to prevent system integrity from application accessing hardware directly.
- Mode Bit is added to the hardware to indicate current mode.
- To enforce this protection, CPU provides two modes of operation:
 - Kernel Mode
 - User Mode

- User Mode

- Direct access to the hardware is prohibited, and so is any arbitrary switching to kernel mode.
- For a user-mode application, Windows creates a process for the application.
- Process provides the application with a private virtual address space and a private handle table.
- As virtual address space is private, one application cannot alter data that belongs to another application.
- If an application crashes, the crash is limited to that one application

- Kernel Mode

- Known as supervisor mode or privileged mode.
- Has complete access to all of the computer's hardware and can control the switching between the CPU modes.
- Code that runs in kernel mode shares a single virtual address space
- If a kernel-mode driver accidentally writes to the wrong virtual address, data that belongs to the operating system or another driver could be compromised
- If a kernel-mode driver crashes, the entire operating system crashes

Operating system classification

Operating Systems can be classified as:


- *GUI*: Graphical User Interface operating systems are operating systems that have the capability of using a mouse and are graphical
- *Multi user*: allows multiple users to utilize the computer and run programs at the same time
- *Multi processing*: allows multiple processors to be utilized
- *Multi tasking*: allows multiple software processes to be run at the same time
- *Multi threading*: allows different parts of a software program to run concurrently

Distributed Operating Systems

- Distributed systems are loosely coupled systems
- A Distributed computer system is a collection of autonomous computer systems
- Distributed systems communicate with one another through various communication lines like high speed buses or telephone lines
- The processors in a distributed system may vary in size and function
- Example: small microprocessors, workstations , minicomputer and large general purpose computers

Distributed Operating Systems

- Processors in distributed systems are referred by no. of different names like, sites, nodes, computers, etc.
- Important reasons for building distributed systems are
 - Resource sharing
 - Computation speedup
 - Reliability
 - Communication
- The key objective of a distributed operating system is transparency
- Ideally, component and resource distribution should be hidden from users and applications programs unless they explicitly demand

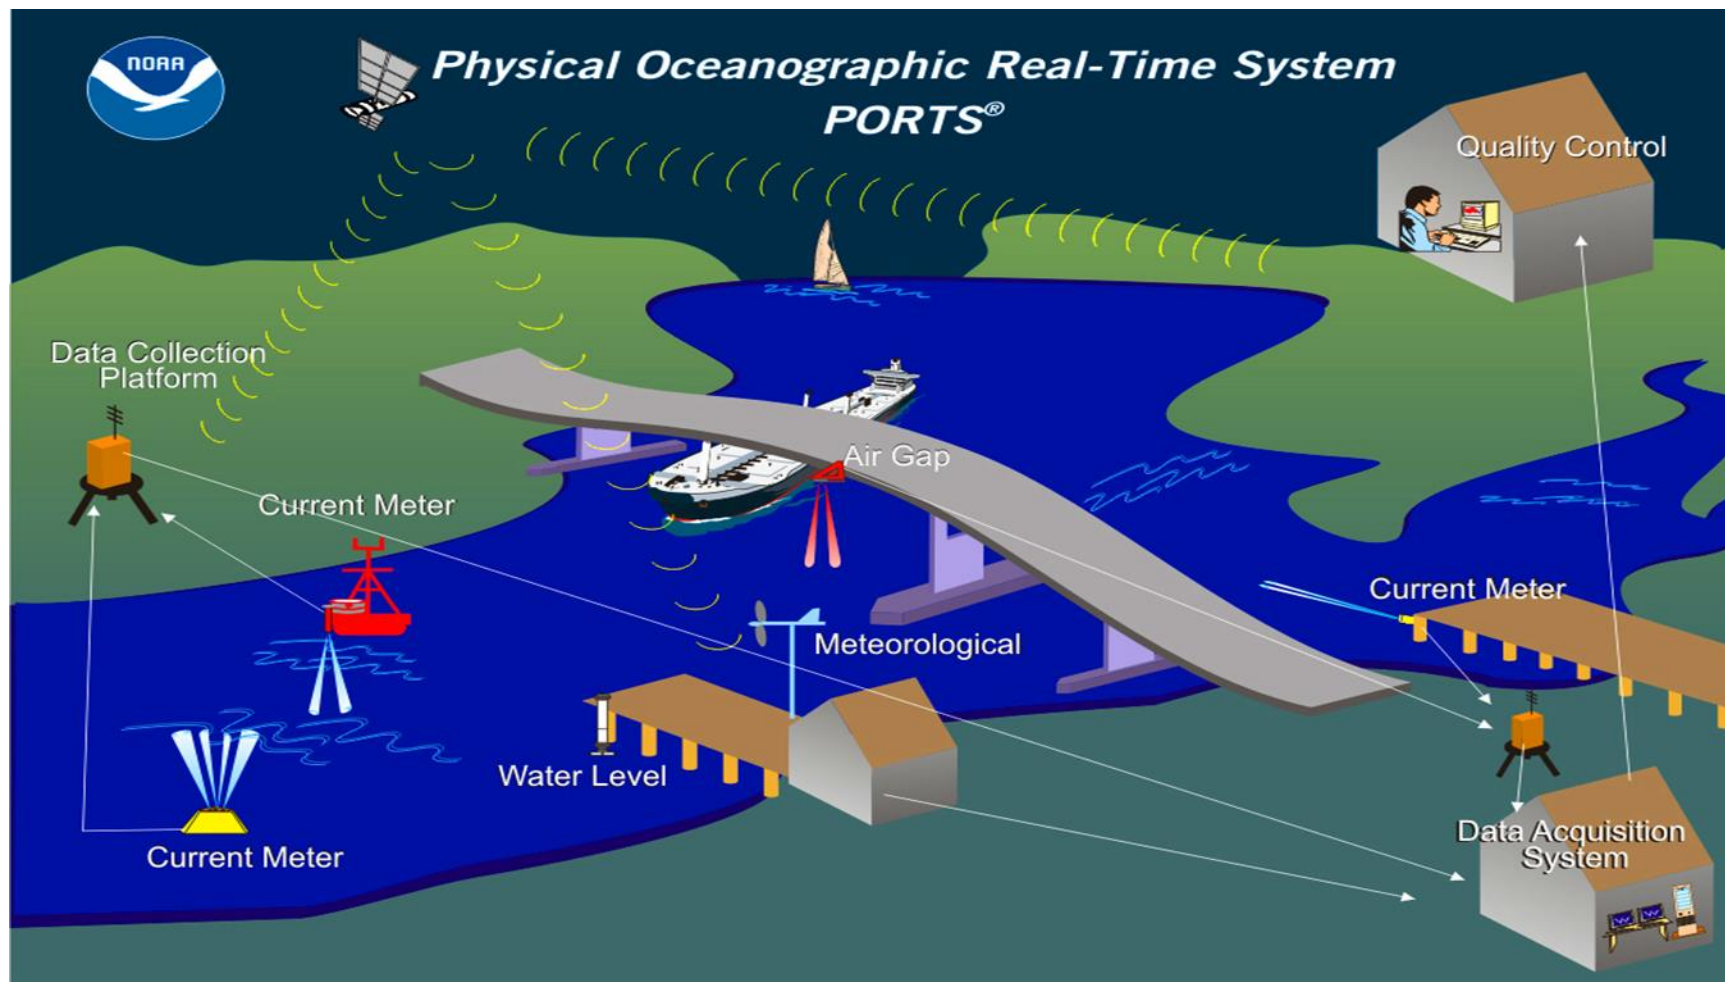


Real Time Operating System

- Real-time systems has well defined, fixed time constraints
- Processing must be done with in the defined constraints, or the system will fail
- For example:
 - Weather Forecasting, Self driving Cars, Earthquake warning systems, wireless sensor networks



Physical Oceanographic Real-Time System PORTS[®]



Real Time Operating System

- Real-time system is used as a control device in a dedicated application
- Sensors bring data to computers
- Computer analyze data and adjust controls to modify the sensor inputs
- Example:
 - Scientific experiments, medical imaging systems, industrial control systems etc..
- Real-time system functions correctly only if it returns the correct result within its time constraints

Real Time Operating System

- A primary objective of real-time systems is to provide quick event - response times, thus meet the scheduling dead lines
- User convenience and resource utilization are of secondary concern to real- time system designers
- Real-time operating systems usually rely on some specific policies and techniques for doing their job

System Calls

- The mechanism used by an application program to request service from the operating system.
- System calls often use a special machine code instruction which causes the processor to change mode (Protected or Supervisor mode)
- This allows the OS to perform restricted actions such as accessing hardware devices or the memory management unit

Types of System Call

Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

Types of System Call

Directory and file system management

Call	Description
s = mkdir(name, mode)	Create a new directory
s = rmdir(name)	Remove an empty directory
s = link(name1, name2)	Create a new entry, name2, pointing to name1
s = unlink(name)	Remove a directory entry
s = mount(special, name, flag)	Mount a file system
s = umount(special)	Unmount a file system

Miscellaneous

Call	Description
s = chdir(dirname)	Change the working directory
s = chmod(name, mode)	Change a file's protection bits
s = kill(pid, signal)	Send a signal to a process
seconds = time(&seconds)	Get the elapsed time since Jan. 1, 1970

System Calls Vs API Call

- Processes in a system runs in different modes, process running in user mode have no access to the privileged instructions.
- If process want perform any privileged instruction or need of any services they request kernel for that service through **System Calls**.
- API is generic term used to identify the functions exposed by any libraries.
- These functions are implemented as part of libraries, or SDK.
- System call is when you call the kernel, whereas a System API are used to invoke system call.


Computing Environments

- Traditional computer
 - Office environment
 - PCs connected to a network, terminals attached to mainframe or minicomputers providing batch and timesharing
 - Now portals allowing networked and remote systems access to same resources
 - Home networks
 - Used to be single system, then modems
 - Now firewalled, networked

Computing Environments

■ Client-Server Computing

- Dumb terminals supplanted by smart PCs
- Many systems now **servers**, responding to requests generated by **clients**
 - ▶ **Compute-server** provides an interface to client to request services (i.e., database)



Peer to Peer Computing

- Another model of distributed system
- P2P does not distinguish clients and servers
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - Registers its service with central lookup service on network, or
 - Broadcast request for service and respond to requests for service via **discovery protocol**
 - Examples include *Napster* and *Torrentz*

Web Based Computing

- Web has become ubiquitous
- PCs most prevalent devices
- More devices becoming networked to allow web access
- New category of devices to manage web traffic among similar servers: **load balancers**
- Use of operating systems like Windows 95, client-side, have evolved into Linux and Windows XP, which can be clients and servers

Thank You